

CARLSON SCHOOL
OF MANAGEMENT
UNIVERSITY OF MINNESOTA


Data Analytics with pandas

1

Carlson School of Management

Outline

- **pandas** Library
- **Series** Object
 - **Index** object
 - Vectorized operations
- **DataFrame** Object
 - Column/row operations
 - Loading data from CSV and Excel files
 - Subsetting with relational and logical operators
 - Splitting with group by and aggregation




2

Carlson School of Management

pandas Library

- Most popular Python data analytics toolkit
 - Built on top of NumPy vectorized operations
- Extremely flexible data structures
 - Reading/writing from/to numerous data sources (CSV, JSON, ...)
 - Splitting, combining, merging, reshaping, joining data
- Fast computations with Series & DataFrames
 - Subsetting, summarizing, pivot tables, ...
 - Tightly integrated with **matplotlib** and **datetime** libraries
- Not part of standard Python installation
 - OS command prompt> **pip install pandas**
 - **LectPD_Cust_Loans.py**
 - >>> import pandas as pd**




3

Carlson School of Management

Fast
Name
0 Series
1 Series
2 Series
3 Series
4 Series
5 Series
6 Series
7 Series
8 Series
9 Series

Index

Value


PANDAS
SERIES

Series Object

- An ordered, 1-dimensional list of data with an index
 - Unlike lists, each **series** element must be of the same data type

```
>>> rates_list = [0.07, 0.075, 0.07, 0.065, 0.077]
>>> rates = pd.Series(rates_list)
```

Displaying the series

```
>>> rates
0    0.070
1    0.075
2    0.070
3    0.065
4    0.077
dtype: float64
```

```
>>> rates.values # NumPy array
>>> rates.values.tolist() # Original list
>>> rates.index # RangeIndex object
>>> rates.index.tolist() # [0, 1, 2, 3, 4]
```


4

Carlson School of Management

Fast
Name
0 Series
1 Series
2 Series
3 Series
4 Series
5 Series
6 Series
7 Series
8 Series
9 Series

Index

Value


PANDAS
SERIES

Index Object

- Index object
 - Used to reference a single or multiple series elements
 - Most often a consecutive list of integers 0, 1, 2, ...
 - Could be any list of unique values – think dictionary keys
- Index derived from a list of loan ID's

```
>>> loan_type_list = ['Mortg', 'Mortg', 'Mortg', 'Car', 'Car']
>>> loan_ids = list(range(1022, 1027))
>>> loan_types = pd.Series(loan_type_list, index=loan_ids)
```
- Series from dictionary → key becomes index

```
>>> amt_dict = {1022: 200000, 1023: 150000,
                  1024: 100000, 1025: 25000, 1026: 10000}
>>> amounts = pd.Series(amt_dict)
```


5

Carlson School of Management

Fast
Name
0 Series
1 Series
2 Series
3 Series
4 Series
5 Series
6 Series
7 Series
8 Series
9 Series

Index

Value


PANDAS
SERIES

Vectorized Operations

- Addition, multiplication, ...
- Index object must be lined up / aligned
 - Cannot multiply rates and amounts as is

```
>>> rates * amounts
```

Redefine rates with loan IDs as indices

```
>>> rates = pd.Series(rates_list, index=loan_ids)
>>> rates * amounts
```

Typically best to leave default indices unless there is a compelling reason to change them

```
>>> rates * amounts
1022    14000.0
1023    11250.0
1024     7000.0
1025     1625.0
1026     770.0
dtype: float64
```

6

2

DataFrame Object

- Two-dimensional table of data
 - Columns represent attributes or characteristics of entities
 - Created from a distinct set of Series objects
 - Rows represent different instances of these entities
 - Both columns and rows indexed
- Loans data frame
 - See **loans.py**
 - Assemble individual series into a list
 - Create data frame with **concat**

```
>>> loans_df
   0      1      2      3      4
0 1022  0.070  Mortg 200000  15
1 1023  0.075  Mortg 150000  15
2 1024  0.070  Mortg 100000  30
3 1025  0.065   Car  25000   3
4 1026  0.077   Car  10000   5
```

```
>>> loan_series = loans.loan_series()
>>> loans_df = pd.concat(loan_series, axis=1)
```

7

DataFrame Function

- Redefine column indices as attribute names

```
>>> loan_cols = ['loanID', 'intRate', ..., 'loanTerm']
```
- Creating data frame from dict of individual series

```
>>> loan_series_dict = dict(zip(loan_cols, loan_series))
```

 - Use **DataFrame** function to create the same data frame object

```
>>> loans_df = pd.DataFrame(loan_series_dict)
```
- Examine data frame contents

```
>>> loans_df.info()
```
- Basic column summary stats

```
>>> loans_df.describe()
```

```
>>> loans_df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   loanID      5 non-null      int64
 1   intRate     5 non-null      float64
 2   loanType    5 non-null      object
 3   amount      5 non-null      int64
 4   loanTerm    5 non-null      int64
dtypes: float64(1), int64(3), object(1)
memory usage: 328.0+ bytes
```

8

Working with DataFrame Columns

- Selecting a single column by name

```
>>> loans_df['intRate']
```
- Selecting multiple columns using names
 - Must assemble column names into a list

```
>>> loans_df[['loanType', 'amount']]
```
- Selecting a column by number
 - loanTerm is the 5th column at index 4

```
>>> loans_df[loans_df.columns[4]]
```
- Selecting multiple columns by slicing
 - First 3 columns (0, 1, and 2)

```
>>> loans_df[loans_df.columns[:3]]
```
- Selecting non-adjacent columns

```
>>> loans_df[loans_df.columns[[1,2,4]]]
```

9

Carlson School of Management

Working with DataFrame Rows

- Use index slicing when possible
 - Select a single row
- Select first 3 rows
- Select the last 2 rows using negative indexing
- Use **loc** when slicing not possible
 - Select non-consecutive rows
- Select rows based on logical (Boolean) column expressions
 - Select all mortgage loans

```
>>> loans_df[1:2]
>>> loans_df[:3]
>>> loans_df[-2:]
>>> loans_df.loc[[1,3,4]]
>>> loans_df[loans_df['loanType']=='Mortg']
```

10

Carlson School of Management

Subsetting with Logical Expressions

- Must review logical and relational operators
- Logical AND operator & (amp symbol)
 - Multiple relational comparisons must be in ()
 - Loans with intRate 7% or higher with amounts over 100K
- Combining multiple relational and logical operators
 - Loans with rates 7% or higher who's terms are either less than 10 or greater than 20 years
- Subsetting rows and columns
 - Mortgage loans, omitting loanType column

```
>>> loans_df[(loans_df['intRate']>=0.07) &
              (loans_df['amount']>100000)]
>>> loans_df[(loans_df['intRate']>=0.07) &
              ((loans_df['loanTerm']<10) |
               (loans_df['loanTerm']>20))]
```

!x	Not x
x y	x OR y
x & y	x AND y

<	less than
<=	less than or equal to
>	greater than
>=	greater than or equal to
==	exactly equal to
!=	not equal to

```
>>> loans_df.loc[loans_df['loanType']=='Mortg',
                  loans_df.columns != 'loanType']
```

11

Carlson School of Management

Reading from CSV and Excel Files


- Read from CSV file into panda's data frame
- Read from Excel file into panda's data frame
 - Need to install **xlrd** library first
- Other sources (not covered in class)
 - Relational Database Management Systems (RDBMS)
 - Java Script Object Notation (JSON) files
 - Scraping Web pages for HTML tables, etc..

```
>>> loans_df = pd.read_csv('Loans.csv')
>>> loans_df.info()
>>> loans_df = pd.read_excel('Loans.xls')
>>> loans_df.head()
>>> loans_df.tail()
```

12

Carlson School of Management

Pandas



More Subsetting Examples

- Easiest column subsetting is to create a list of quoted names
 - Mortgage loans showing loanID, amount, rate, term and payment

```
>>> loans_df.loc[loans_df['loanType']=='Mortg',  
                ['loanID', 'amount', 'intRate', 'loanTerm', 'mthPmt']]
```
- More complex column subsetting with exclusion operator ~ (tilda symbol) and **isin** function
 - Mortgage loans excluding first and last name and loan type


```
>>> loans_df.loc[loans_df['loanType']=='Mortg',  
                ~loans_df.columns.isin(['firstName', 'lastName', 'loanType'])]
```
- The use of **!=** relational operator
 - Showing all loans from customers that are not from Taos

```
>>> loans_no_Taos_df = loans_df.loc[loans_df['city']!='Taos', ...]
```

13

Carlson School of Management

Pandas



More Subsetting Examples (cont.)

- Using relational and logical operators in complex expressions
 - Show all Taos mortgage loans

```
>>> loans_Taos_Mortg_df = loans_df.loc[(loans_df['city']=='Taos')  
    & (loans_df['loanType']=='Mortg'),...]
```
- Working with both AND (&) and OR (|) logical operators
 - Mortgage loans either over half a million or under 200K


```
>>> loans_mortg_high_low_df = loans_df.loc[(loans_df['loanType']=='Mortg')  
    & ((loans_df['amount']>500000) | (loans_df['amount']<200000)),...]
```
- Working with datetimes
 - Issues with **loanDate** column
 - NOT recognized as **datetime** from CSV file
 - IS recognized as **datetime** from Excel file
 - Showing all January loans (see **LectPD_Cust_Loans.py** for details)

```
>>> loans_jan_df = loans_df.loc[(loans_df['loanDate']>=beg_jan) &  
    (loans_df['loanDate']<=end_jan), ...]
```

14

Carlson School of Management

Pandas



Sorting and Adding Columns

- Sorting on one column
 - Provide the column to sort on and the sort order

```
>>> loans_sort1_df = loans_sub_df.sort_values('mthPmt',  
    ascending=False)
```
- Sorting on multiple columns
 - Create a list of columns, first one must be categorical and sort order tuple

```
>>> loans_sort2_df = loans_sub_df.sort_values(['loanType',  
    'amount'], ascending=(True, False))
```
- Creating a new column
 - Provide a name and an expression involving existing columns

```
>>> loans_df['totPmt'] = loans_df['mthPmt'] *  
    loans_df['loanTerm'] * 12
```

15

Carlson School of Management

Pandas

Group By (Splitting) and Aggregating

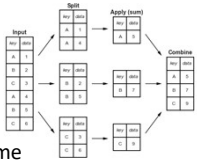
- Split -> Apply (Aggregate) -> Combine
- Split loans into 3 groups by loan type
 - Mortgage, Car and Other loans

```
>>> loans_df.groupby('loanType')
```
- Apply aggregate operation on each group and combine the result into a new data frame
 - Average monthly payment

```
>>> loans_df.groupby('loanType').mean()['mthPmt']
```
 - Total amount borrowed

```
>>> loans_df.groupby('loanType').sum()['amount']
```
 - Number of loans by type

```
>>> loans_df.groupby('loanType').count()['loanID']
```



16

Carlson School of Management

Pandas

Grouping By Multiple Columns

- Combine two (or more) categorical variables
 - Find the average monthly payment by loan type and city

```
>>> loans_df.groupby(['loanType', 'city']).mean()['mthPmt']
```
- Use multiple aggregations
 - Find the min, max and average monthly payment by loan type
 - One element dictionary with column name as a key and aggregate operations as a list of values

```
>>> loans_df.groupby('loanType').agg({'mthPmt': ['min', 'max', 'mean']})
```
- Group by multiple columns and perform multiple aggregations
 - Find the number of loans; total and average of amounts borrowed; min, max and average of monthly payments by loan type and city
 - Three element dictionary with three column names as keys, followed by lists of operations on those columns as values


```
>>> agg_dict = {'loanID': ['count'], 'amount': ['sum', 'mean'], 'mthPmt': ['min', 'max', 'mean']}
>>> loans_df.groupby(['loanType', 'city']).agg(agg_dict)
```

17

Carlson School of Management

Summary

- Introduced **pandas** library for data analysis
- Defined **Series** object
 - Supports vectorized operations using **Index** object
- Defined **DataFrame** object
 - The most important data structure for doing data analytics with pandas
 - Columns consists of series objects; rows represent different observations (instances) of various entities
 - Demonstrated how to work with columns and rows
- Loaded data frames from CSV and Excel files
 - Presented a variety of **subsetting** and **summarizing** operations with **relational** and **logical** operators
 - Showed how to **sort** on one or more columns
 - Finished with Split-Apply-Combine operations using **groupby** and aggregation functions



18